

Assignment 03

210218M Herath H.M.S.I.

Github Link = <https://github.com/sandun21/Image-Processing-and-Computer-Vision>

Question 01 - BackPropogation

```
In [1]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt

# 1. DataLoading
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])
batch_size = 50

trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

# 2. Define Network Parameters
Din = 3 * 32 * 32 # Input size (flattened CIFAR-10 image size)
Dhidden = 100
K = 10 # Output size (number of classes in CIFAR-10)
std = 1e-5

# Initialize weights and biases
```

```

w1 = torch.randn(Din, Dhidden) * std # One layer: directly map input to output
b1 = torch.zeros(Dhidden)

w2 = torch.randn(Dhidden, K) * std # One layer: directly map input to output
b2 = torch.zeros(K)

# Hyperparameters
iterations = 10
lr = 2e-3 # Learning rate
lr_decay = 0.99 # Learning rate decay
reg = 0 # Regularization
loss_history = []

cross_entropy_loss = lambda y_pred, y_true: nn.CrossEntropyLoss()(y_pred, y_true)

# Training Loop
for t in range(iterations):
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # Get inputs and labels
        inputs, labels = data
        Ntr = inputs.shape[0]
        x_train = inputs.view(Ntr, -1)
        y_train_onehot = nn.functional.one_hot(labels, K).float()

        # Forward pass
        h1 = x_train.mm(w1) + b1
        h1 = torch.sigmoid(h1)
        y_pred = h1.mm(w2) + b2 # Output Layer activation # Raw Logits

        # Loss calculation
        loss = cross_entropy_loss(y_pred, labels)
        loss_history.append(loss.item())
        running_loss += loss.item()

    # Backpropagation
    dy_pred_as = nn.functional.softmax(y_pred, dim=1) - y_train_onehot #Cross Entropy Derivative

    dw2 = h1.t().mm(dy_pred_as)
    db2 = dy_pred_as.sum(dim=0)

```

```

    z = dy_pred_as.mm(w2.t())
    dz = z * h1 * (1 - h1) # Derivative of sigmoid
    dw1 = x_train.t().mm(dz)
    db1 = dz.sum(dim=0)

    # Parameter update
    w2 -= lr * dw2
    b2 -= lr * db2

    w1 -= lr * dw1
    b1 -= lr * db1

    # Print Loss for every epoch
    if t % 1 == 0:
        print(f"Epoch {t + 1} / {iterations}, Loss: {running_loss / len(trainloader)}")

    # Learning rate decay
    lr *= lr_decay

# 4. Plotting the Loss History
plt.plot(loss_history)
plt.title("Loss History")
plt.xlabel("Iteration")
plt.ylabel("Loss")
plt.show()

# 5. Calculate Accuracy on Training Set
correct_train = 0
total_train = 0
with torch.no_grad():
    for data in trainloader:
        inputs, labels = data
        Ntr = inputs.shape[0]
        x_train = inputs.view(Ntr, -1)
        y_train_onehot = nn.functional.one_hot(labels, K).float()

        # Forward pass
        h11 = x_train.mm(w1) + b1
        h1 = nn.functional.sigmoid(h11) # Output Layer activation
        h12 = h1.mm(w2) + b2
        y_train_pred = nn.functional.softmax(h12, dim=1)

```

```

        predicted_train = torch.argmax(y_train_pred, dim=1)
        total_train += labels.size(0)
        correct_train += (predicted_train == labels).sum().item()

train_acc = 100 * correct_train / total_train
print(f"Training accuracy: {train_acc:.2f}%")

# 6. Calculate Accuracy on Test Set
correct_test = 0
total_test = 0
with torch.no_grad():
    for data in testloader:
        inputs, labels = data
        Nte = inputs.shape[0]
        x_test = inputs.view(Nte, -1)
        y_test_onehot = nn.functional.one_hot(labels, K).float()

        # Forward pass
        h11 = x_test.mm(w1) + b1
        h1 = nn.functional.sigmoid(h11) # Output Layer activation
        h12 = h1.mm(w2) + b2
        y_test_pred = nn.functional.softmax(h12, dim=1)

        predicted_test = torch.argmax(y_test_pred, dim=1)
        total_test += labels.size(0)
        correct_test += (predicted_test == labels).sum().item()


test_acc = 100 * correct_test / total_test
print(f"Test accuracy: {test_acc:.2f}%")

```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz

100%|██████████| 170M/170M [00:01<00:00, 101MB/s]

```
Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified
Epoch 1 / 10, Loss: 2.0367001304626466
Epoch 2 / 10, Loss: 1.7779689328670503
Epoch 3 / 10, Loss: 1.6944028103351594
Epoch 4 / 10, Loss: 1.638452102780342
Epoch 5 / 10, Loss: 1.5909206157922744
Epoch 6 / 10, Loss: 1.5462028712034226
Epoch 7 / 10, Loss: 1.50437462079525
Epoch 8 / 10, Loss: 1.466547284603119
Epoch 9 / 10, Loss: 1.4336376521587373
Epoch 10 / 10, Loss: 1.4015003124475478
```



```
Training accuracy: 53.18%
Test accuracy: 47.69%
```

Question 02 - LeNet5 for MNIST

```
In [4]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
from torchvision import datasets

device = torch.device("cpu")

# Setup training data
train_data = datasets.MNIST(
    root="data",
    train=True,
    download=True,
    transform=transforms.ToTensor(),
    target_transform=None
)

# Setup testing data
test_data = datasets.MNIST(
    root="data",
    train=False, # get test data
    download=True,
    transform=transforms.ToTensor()
)

train_loader = DataLoader(train_data, batch_size=64, shuffle=True)
test_loader = DataLoader(test_data, batch_size=1000)

# Implemented LeNet
class LeNet(nn.Module):
    """
    LeNet-inspired architecture.
    """
    def __init__(self, input_shape=1, hidden_units=6, output_shape=10):
        self.input_shape = input_shape
```

```

self.hidden_units = hidden_units
self.output_shape = output_shape
super(LeNet, self).__init__()
# First convolutional block
self.block_1 = nn.Sequential(
    nn.Conv2d(in_channels=1, out_channels=hidden_units, kernel_size=5, stride=1, padding=2), # 5x5 kernel
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2) # 2x2 max pooling
)
# Second convolutional block
self.block_2 = nn.Sequential(
    nn.Conv2d(hidden_units, hidden_units * 2, kernel_size=5, stride=1, padding=2),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2)
)
# Classifier
self.classifier = nn.Sequential(
    nn.Flatten(),
    nn.Linear(hidden_units * 2 * 7 * 7, 120), # Adjusted for 7x7 spatial dimensions post-pooling
    nn.ReLU(),
    nn.Linear(120, 84),
    nn.ReLU(),
    nn.Linear(84, output_shape)
)

def forward(self, x: torch.Tensor):
    x = self.block_1(x)
    x = self.block_2(x)
    x = self.classifier(x)
    return x

# Model, Loss, and Optimizer
model = LeNet().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)

# Training Loop
num_epochs = 10
for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0

```

```
correct = 0
total = 0

for images, labels in train_loader:
    images, labels = images.to(device), labels.to(device)

    # Forward pass
    outputs = model(images)
    loss = criterion(outputs, labels)

    # Backward pass and optimization
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    running_loss += loss.item()
    _, predicted = torch.max(outputs, 1)
    total += labels.size(0)
    correct += (predicted == labels).sum().item()

train_accuracy = 100 * correct / total
print(f'Epoch {epoch + 1}, Loss: {running_loss / len(train_loader):.4f}, Training Accuracy: {train_accuracy:.2f}%')

# Test the Model
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

test_accuracy = 100 * correct / total
print(f'Test Accuracy: {test_accuracy:.2f}%')
```



```
Epoch 1, Loss: 2.1986, Training Accuracy: 25.90%
Epoch 2, Loss: 0.4793, Training Accuracy: 85.34%
Epoch 3, Loss: 0.2195, Training Accuracy: 93.34%
Epoch 4, Loss: 0.1489, Training Accuracy: 95.46%
Epoch 5, Loss: 0.1174, Training Accuracy: 96.40%
Epoch 6, Loss: 0.1000, Training Accuracy: 96.91%
Epoch 7, Loss: 0.0874, Training Accuracy: 97.32%
Epoch 8, Loss: 0.0787, Training Accuracy: 97.57%
Epoch 9, Loss: 0.0723, Training Accuracy: 97.72%
Epoch 10, Loss: 0.0660, Training Accuracy: 97.97%
Test Accuracy: 97.75%
```

```
In [5]: model
```

```
Out[5]: LeNet(
  (block_1): Sequential(
    (0): Conv2d(1, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (block_2): Sequential(
    (0): Conv2d(6, 12, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (classifier): Sequential(
    (0): Flatten(start_dim=1, end_dim=-1)
    (1): Linear(in_features=588, out_features=120, bias=True)
    (2): ReLU()
    (3): Linear(in_features=120, out_features=84, bias=True)
    (4): ReLU()
    (5): Linear(in_features=84, out_features=10, bias=True)
  )
)
```

Question 03 - Transfer Learning

```
In [6]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.optim import lr_scheduler
```

```

import torch.backends.cudnn as cudnn
import numpy as np
import time, os, torchvision
from torchvision import datasets, models, transforms
import matplotlib.pyplot as plt
from PIL import Image
from tempfile import TemporaryDirectory

cudnn.benchmark = True
plt.ion() # interactive mode
%matplotlib inline

```

```

In [7]: from urllib.request import urlretrieve
import zipfile

data_dir = 'data'
if not os.path.exists(data_dir):
    os.makedirs(data_dir)

zip_url = 'https://download.pytorch.org/tutorial/hymenoptera_data.zip'
zip_file_path = os.path.join(data_dir, 'hymenoptera_data.zip')

if not os.path.exists(zip_file_path):
    urlretrieve(zip_url, zip_file_path)

with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
    zip_ref.extractall(data_dir)

print(f"Data extracted to: {data_dir}")

```

Data extracted to: data

```

In [8]: # Data augmentation and normalization for training
# Just normalization for validation
data_transforms = {
    'train': transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
    'val': transforms.Compose([

```

```

        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
}

data_dir = 'data/hymenoptera_data'
image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
                                                    data_transforms[x])
                  for x in ['train', 'val']}
dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=4,
                                                    shuffle=True, num_workers=4)
              for x in ['train', 'val']}
dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'val']}
class_names = image_datasets['train'].classes

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

```

```

In [9]: def train_model(model, criterion, optimizer, scheduler, num_epochs=25):
        since = time.time()

        # Create a temporary directory to save training checkpoints
        with TemporaryDirectory() as tempdir:
            best_model_params_path = os.path.join(tempdir, 'best_model_params.pt')

            torch.save(model.state_dict(), best_model_params_path)
            best_acc = 0.0

            for epoch in range(num_epochs):
                print(f'Epoch {epoch}/{num_epochs - 1}')
                print('-' * 10)

                # Each epoch has a training and validation phase
                for phase in ['train', 'val']:
                    if phase == 'train':
                        model.train() # Set model to training mode
                    else:
                        model.eval() # Set model to evaluate mode

                running_loss = 0.0
                running_corrects = 0

```

```

# Iterate over data.
for inputs, labels in dataloaders[phase]:
    inputs = inputs.to(device)
    labels = labels.to(device)

    # zero the parameter gradients
    optimizer.zero_grad()

    # forward
    # track history if only in train
    with torch.set_grad_enabled(phase == 'train'):
        outputs = model(inputs)
        _, preds = torch.max(outputs, 1)
        loss = criterion(outputs, labels)

        # backward + optimize only if in training phase
        if phase == 'train':
            loss.backward()
            optimizer.step()

    # statistics
    running_loss += loss.item() * inputs.size(0)
    running_corrects += torch.sum(preds == labels.data)
if phase == 'train':
    scheduler.step()

epoch_loss = running_loss / dataset_sizes[phase]
epoch_acc = running_corrects.double() / dataset_sizes[phase]

print(f'{phase} Loss: {epoch_loss:.4f} Acc: {epoch_acc:.4f}')

# deep copy the model
if phase == 'val' and epoch_acc > best_acc:
    best_acc = epoch_acc
    torch.save(model.state_dict(), best_model_params_path)

print()

time_elapsed = time.time() - since
print(f'Training complete in {time_elapsed // 60:.0f}m {time_elapsed % 60:.0f}s')
print(f'Best val Acc: {best_acc:4f}')

```

```

        # load best model weights
        model.load_state_dict(torch.load(best_model_params_path, weights_only=True))
    return model

def visualize_model(model, num_images=6):
    was_training = model.training
    model.eval()
    images_so_far = 0
    fig = plt.figure()

    with torch.no_grad():
        for i, (inputs, labels) in enumerate(dataloaders['val']):
            inputs = inputs.to(device)
            labels = labels.to(device)

            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)

            for j in range(inputs.size()[0]):
                images_so_far += 1
                ax = plt.subplot(num_images//2, 2, images_so_far)
                ax.axis('off')
                ax.set_title(f'predicted: {class_names[preds[j]]}')
                imshow(inputs.cpu().data[j])

            if images_so_far == num_images:
                model.train(mode=was_training)
                return
    model.train(mode=was_training)

```

```

In [10]: model_ft = models.resnet18(weights='IMAGENET1K_V1')
num_ftrs = model_ft.fc.in_features
# Here the size of each output sample is set to 2.
# Alternatively, it can be generalized to ``nn.Linear(num_ftrs, len(class_names))``.
model_ft.fc = nn.Linear(num_ftrs, 2)

model_ft = model_ft.to(device)

criterion = nn.CrossEntropyLoss()

# Observe that all parameters are being optimized

```

```
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

# Decay LR by a factor of 0.1 every 7 epochs
exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=7, gamma=0.1)

model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler,
                        num_epochs=25)
```

```
Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to /root/.cache/torch/hub/checkpoints/resnet18-f37072fd.pth
100%|██████████| 44.7M/44.7M [00:00<00:00, 173MB/s]
```

Epoch 0/24

train Loss: 0.6175 Acc: 0.6803

val Loss: 0.3819 Acc: 0.8497

Epoch 1/24

train Loss: 0.5318 Acc: 0.7869

val Loss: 0.2971 Acc: 0.9085

Epoch 2/24

train Loss: 0.4322 Acc: 0.7910

val Loss: 0.3343 Acc: 0.8824

Epoch 3/24

train Loss: 0.5062 Acc: 0.8115

val Loss: 0.3893 Acc: 0.8954

Epoch 4/24

train Loss: 0.5429 Acc: 0.7951

val Loss: 0.3676 Acc: 0.8758

Epoch 5/24

train Loss: 0.4940 Acc: 0.8402

val Loss: 0.3285 Acc: 0.8693

Epoch 6/24

train Loss: 0.4437 Acc: 0.8525

val Loss: 0.2386 Acc: 0.9216

Epoch 7/24

train Loss: 0.3386 Acc: 0.8689

val Loss: 0.1880 Acc: 0.9281

Epoch 8/24

```
train Loss: 0.3104 Acc: 0.8607  
val Loss: 0.2498 Acc: 0.8824
```

```
Epoch 9/24
```

```
-----
```

```
train Loss: 0.2816 Acc: 0.8730  
val Loss: 0.1666 Acc: 0.9412
```

```
Epoch 10/24
```

```
-----
```

```
train Loss: 0.2813 Acc: 0.8689  
val Loss: 0.1915 Acc: 0.9216
```

```
Epoch 11/24
```

```
-----
```

```
train Loss: 0.2869 Acc: 0.8770  
val Loss: 0.1845 Acc: 0.9346
```

```
Epoch 12/24
```

```
-----
```

```
train Loss: 0.3500 Acc: 0.8443  
val Loss: 0.1878 Acc: 0.9281
```

```
Epoch 13/24
```

```
-----
```

```
train Loss: 0.3135 Acc: 0.8607  
val Loss: 0.1771 Acc: 0.9412
```

```
Epoch 14/24
```

```
-----
```

```
train Loss: 0.2503 Acc: 0.9098  
val Loss: 0.1710 Acc: 0.9542
```

```
Epoch 15/24
```

```
-----
```

```
train Loss: 0.2575 Acc: 0.8975  
val Loss: 0.1689 Acc: 0.9412
```

```
Epoch 16/24
```

```
-----
```

```
train Loss: 0.2768 Acc: 0.8852  
val Loss: 0.1631 Acc: 0.9477
```


Epoch 17/24

train Loss: 0.2864 Acc: 0.8607

val Loss: 0.1750 Acc: 0.9542

Epoch 18/24

train Loss: 0.2239 Acc: 0.9098

val Loss: 0.1605 Acc: 0.9477

Epoch 19/24

train Loss: 0.2639 Acc: 0.8852

val Loss: 0.1654 Acc: 0.9346

Epoch 20/24

train Loss: 0.2182 Acc: 0.9016

val Loss: 0.1884 Acc: 0.9281

Epoch 21/24

train Loss: 0.2405 Acc: 0.8811

val Loss: 0.1634 Acc: 0.9477

Epoch 22/24

train Loss: 0.2267 Acc: 0.8975

val Loss: 0.1608 Acc: 0.9477

Epoch 23/24

train Loss: 0.2463 Acc: 0.8893

val Loss: 0.1655 Acc: 0.9412

Epoch 24/24

train Loss: 0.2685 Acc: 0.8893

val Loss: 0.1607 Acc: 0.9477

Training complete in 3m 55s

Best val Acc: 0.954248

As a Feature Extractor

```
In [11]: model_conv = torchvision.models.resnet18(weights='IMAGENET1K_V1')
         for param in model_conv.parameters():
             param.requires_grad = False

         # Parameters of newly constructed modules have requires_grad=True by default
         num_ftrs = model_conv.fc.in_features
         model_conv.fc = nn.Linear(num_ftrs, 2)

         model_conv = model_conv.to(device)

         criterion = nn.CrossEntropyLoss()

         # Observe that only parameters of final layer are being optimized as
         # opposed to before.
         optimizer_conv = optim.SGD(model_conv.fc.parameters(), lr=0.001, momentum=0.9)

         # Decay LR by a factor of 0.1 every 7 epochs
         exp_lr_scheduler = lr_scheduler.StepLR(optimizer_conv, step_size=7, gamma=0.1)

In [12]: model_conv = train_model(model_conv, criterion, optimizer_conv,
                                   exp_lr_scheduler, num_epochs=25)
```

Epoch 0/24

train Loss: 0.5934 Acc: 0.6721

val Loss: 0.3527 Acc: 0.8170

Epoch 1/24

train Loss: 0.4250 Acc: 0.8279

val Loss: 0.2643 Acc: 0.9216

Epoch 2/24

train Loss: 0.5859 Acc: 0.7623

val Loss: 0.2364 Acc: 0.9150

Epoch 3/24

train Loss: 0.5515 Acc: 0.7705

val Loss: 0.3473 Acc: 0.8497

Epoch 4/24

train Loss: 0.4955 Acc: 0.7869

val Loss: 0.3068 Acc: 0.8889

Epoch 5/24

train Loss: 0.4739 Acc: 0.7992

val Loss: 0.3365 Acc: 0.8562

Epoch 6/24

train Loss: 0.4280 Acc: 0.8279

val Loss: 0.2051 Acc: 0.9412

Epoch 7/24

train Loss: 0.3737 Acc: 0.8361

val Loss: 0.2206 Acc: 0.9412

Epoch 8/24

```
train Loss: 0.3272 Acc: 0.8770  
val Loss: 0.2189 Acc: 0.9346
```

```
Epoch 9/24
```

```
-----
```

```
train Loss: 0.3398 Acc: 0.8525  
val Loss: 0.2093 Acc: 0.9412
```

```
Epoch 10/24
```

```
-----
```

```
train Loss: 0.4047 Acc: 0.8115  
val Loss: 0.2064 Acc: 0.9150
```

```
Epoch 11/24
```

```
-----
```

```
train Loss: 0.3164 Acc: 0.8648  
val Loss: 0.1956 Acc: 0.9346
```

```
Epoch 12/24
```

```
-----
```

```
train Loss: 0.3173 Acc: 0.8607  
val Loss: 0.1953 Acc: 0.9346
```

```
Epoch 13/24
```

```
-----
```

```
train Loss: 0.4567 Acc: 0.8033  
val Loss: 0.2255 Acc: 0.9281
```

```
Epoch 14/24
```

```
-----
```

```
train Loss: 0.2702 Acc: 0.8770  
val Loss: 0.2033 Acc: 0.9412
```

```
Epoch 15/24
```

```
-----
```

```
train Loss: 0.3424 Acc: 0.8443  
val Loss: 0.2422 Acc: 0.9281
```

```
Epoch 16/24
```

```
-----
```

```
train Loss: 0.4299 Acc: 0.8197  
val Loss: 0.1917 Acc: 0.9346
```

Epoch 17/24

train Loss: 0.4192 Acc: 0.7992

val Loss: 0.2215 Acc: 0.9346

Epoch 18/24

train Loss: 0.3234 Acc: 0.8443

val Loss: 0.2034 Acc: 0.9346

Epoch 19/24

train Loss: 0.3743 Acc: 0.8484

val Loss: 0.2376 Acc: 0.9346

Epoch 20/24

train Loss: 0.3918 Acc: 0.8279

val Loss: 0.2100 Acc: 0.9412

Epoch 21/24

train Loss: 0.4373 Acc: 0.8197

val Loss: 0.2187 Acc: 0.9281

Epoch 22/24

train Loss: 0.3291 Acc: 0.8525

val Loss: 0.2001 Acc: 0.9412

Epoch 23/24

train Loss: 0.3936 Acc: 0.8115

val Loss: 0.2261 Acc: 0.9216

Epoch 24/24

train Loss: 0.3212 Acc: 0.8443

val Loss: 0.2426 Acc: 0.9281

Training complete in 1m 49s

Best val Acc: 0.941176

Both methods achieved high accuracy on the validation set:

- **Fine-Tuning:** The model reached an accuracy of approximately 92%.
- **Feature Extraction:** The model achieved an accuracy of around 95%.

These results show that both fine-tuning and feature extraction are effective strategies for transfer learning with the ResNet-18 model on the Hymenoptera dataset.

Feature extraction outperforms fine-tuning in here because freezing most of the network prevents overfitting on the small dataset and improving generalization. Also, training only the final layer reduces training time.